

Lesson_1

March 5, 2025

1 Lesson 1: Why Pretraining?

1.1 1. Install dependencies and fix seed

Welcome to Lesson 1!

If you would like to access the `requirements.txt` file for this course, go to **File** and click on **Open**.

```
[ ]: # Install any packages if it does not exist
# !pip install -q -r ../requirements.txt
```

```
[ ]: # Ignore insignificant warnings (ex: deprecations)
import warnings
warnings.filterwarnings('ignore')
```

```
[ ]: # Set a seed for reproducibility
import torch

def fix_torch_seed(seed=42):
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

fix_torch_seed()
```

1.2 2. Load a general pretrained model

This course will work with small models that fit within the memory of the learning platform. TinySolar-248m-4k is a small decoder-only model with 248M parameters (similar in scale to GPT2) and a 4096 token context window. You can find the model on the Hugging Face model library at [this link](#).

You'll load the model in three steps: 1. Specify the path to the model in the Hugging Face model library 2. Load the model using `AutoModelForCausalLM` in the `transformers` library 3. Load the tokenizer for the model from the same model path

```
[ ]: model_path_or_name = "../models/TinySolar-248m-4k"
```

```
[ ]: from transformers import AutoModelForCausalLM
tiny_general_model = AutoModelForCausalLM.from_pretrained(
    model_path_or_name,
    device_map="cpu", # change to auto if you have access to a GPU
    torch_dtype=torch.bfloat16
)
```

```
[ ]: from transformers import AutoTokenizer
tiny_general_tokenizer = AutoTokenizer.from_pretrained(
    model_path_or_name
)
```

1.3 3. Generate text samples

Here you'll try generating some text with the model. You'll set a prompt, instantiate a text streamer, and then have the model complete the prompt:

```
[ ]: prompt = "I am an engineer. I love"
```

```
[ ]: inputs = tiny_general_tokenizer(prompt, return_tensors="pt")
```

```
[ ]: from transformers import TextStreamer
streamer = TextStreamer(
    tiny_general_tokenizer,
    skip_prompt=True, # If you set to false, the model will first return the
    ↪ prompt and then the generated text
    skip_special_tokens=True
)
```

```
[ ]: outputs = tiny_general_model.generate(
    **inputs,
    streamer=streamer,
    use_cache=True,
    max_new_tokens=128,
    do_sample=False,
    temperature=0.0,
    repetition_penalty=1.1
)
```

1.4 4. Generate Python samples with pretrained general model

Use the model to write a python function called `find_max()` that finds the maximum value in a list of numbers:

```
[ ]: prompt = "def find_max(numbers):"
```

```
[ ]: inputs = tiny_general_tokenizer(
    prompt, return_tensors="pt"
```

```

).to(tiny_general_model.device)

streamer = TextStreamer(
    tiny_general_tokenizer,
    skip_prompt=True, # Set to false to include the prompt in the output
    skip_special_tokens=True
)

```

```

[ ]: outputs = tiny_general_model.generate(
    **inputs,
    streamer=streamer,
    use_cache=True,
    max_new_tokens=128,
    do_sample=False,
    temperature=0.0,
    repetition_penalty=1.1
)

```

1.5 5. Generate Python samples with finetuned Python model

This model has been fine-tuned on instruction code examples. You can find the model and information about the fine-tuning datasets on the Hugging Face model library at [this link](#).

You'll follow the same steps as above to load the model and use it to generate text.

```

[ ]: model_path_or_name = "./models/TinySolar-248m-4k-code-instruct"

```

```

[ ]: tiny_finetuned_model = AutoModelForCausalLM.from_pretrained(
    model_path_or_name,
    device_map="cpu",
    torch_dtype=torch.bfloat16,
)

tiny_finetuned_tokenizer = AutoTokenizer.from_pretrained(
    model_path_or_name
)

```

```

[ ]: prompt = "def find_max(numbers):"

inputs = tiny_finetuned_tokenizer(
    prompt, return_tensors="pt"
).to(tiny_finetuned_model.device)

streamer = TextStreamer(
    tiny_finetuned_tokenizer,
    skip_prompt=True,
    skip_special_tokens=True
)

```

```

outputs = tiny_finetuned_model.generate(
    **inputs,
    streamer=streamer,
    use_cache=True,
    max_new_tokens=128,
    do_sample=False,
    temperature=0.0,
    repetition_penalty=1.1
)

```

1.6 6. Generate Python samples with pretrained Python model

Here you'll use a version of TinySolar-248m-4k that has been further pretrained (a process called **continued pretraining**) on a large selection of python code samples. You can find the model on Hugging Face at [this link](#).

You'll follow the same steps as above to load the model and use it to generate text.

```
[ ]: model_path_or_name = "./models/TinySolar-248m-4k-py"
```

```
[ ]: tiny_custom_model = AutoModelForCausalLM.from_pretrained(
    model_path_or_name,
    device_map="cpu",
    torch_dtype=torch.bfloat16,
)

tiny_custom_tokenizer = AutoTokenizer.from_pretrained(
    model_path_or_name
)

```

```
[ ]: prompt = "def find_max(numbers):"

inputs = tiny_custom_tokenizer(
    prompt, return_tensors="pt"
).to(tiny_custom_model.device)

streamer = TextStreamer(
    tiny_custom_tokenizer,
    skip_prompt=True,
    skip_special_tokens=True
)

outputs = tiny_custom_model.generate(
    **inputs, streamer=streamer,
    use_cache=True,
    max_new_tokens=128,
    do_sample=False,

```

```
    repetition_penalty=1.1  
)
```

Try running the python code the model generated above:

```
[ ]: def find_max(numbers):  
    max = 0  
    for num in numbers:  
        if num > max:  
            max = num  
    return max
```

```
[ ]: find_max([1,3,5,1,6,7,2])
```

Lesson_2

March 5, 2025

1 Lecture 2: Data Preparation

In this lesson you'll carry out some of the data cleaning steps required to prepare data for pretraining. In the video, Sung mentioned an Upstage tool called **Dataverse** which can help you with data cleaning. You can checkout the features of Dataverse at [this link](#).

```
[ ]: import warnings
      warnings.filterwarnings("ignore")
```

1.1 1. Sourcing datasets for pretraining

In this section, you'll see two ways to source data for training: 1. Download an existing dataset from Hugging Face 2. Create a dataset of python scripts sourced from Github

In both cases the result will be a Hugging Face **Dataset** object, part of the **Datasets** library. You can read more about the properties of Datasets and how to work with them on the [Hugging Face website](#).

1.1.1 Download data from Hugging face

The dataset you download here is a subset of a much larger dataset called **Red Pajama**. The full, 1 trillion token dataset is available on Hugging Face at [this link](#).

```
[ ]: import datasets
      pretraining_dataset = datasets.load_dataset(
          "upstage/Pretraining_Dataset",
          split="train"
      )
```

```
[ ]: print(pretraining_dataset)
```

Only work with the `text` column:

```
[ ]: pretraining_dataset = pretraining_dataset.select_columns(
      ["text"]
  )
```

Print a sample:

```
[ ]: print(pretraining_dataset[0]["text"][:500])
```

1.1.2 Compare pretraining and fine-tuning datasets

In the next cell, you'll download a fine-tuning dataset to contrast with the pretraining dataset you loaded above. You can read more about the Alpaca model and instruction tuning dataset [here](#).

```
[ ]: instruction_dataset = datasets.load_dataset(  
    "c-s-ale/alpaca-gpt4-data",  
    split='train'  
)  
print(instruction_dataset)  
  
[ ]: i=0  
print("Instruction: " + instruction_dataset[i]["instruction"]  
    + "\nInput: " + instruction_dataset[i]["input"]  
    + "\nOutput: " + instruction_dataset[i]["output"])
```

Notice how in contrast to the pretraining data, which is just raw text, fine-tuning datasets are structured into question-answer pairs or instruction-response sets that can include additional input context if required.

Moving forward, you'll only work with the unstructured pretraining dataset.

1.1.3 Scrape python code from Github

Here, you'll download a selection of python scripts from Github and then prepare them as a Hugging Face Dataset object to use in training.

The same pattern here will work for preparing any text scraped from the web.

```
[ ]: # Import some required packages  
import os  
import requests  
  
# Path to directory to store python scripts  
code_dir = "./code"  
  
[ ]: urls = [  
    "https://raw.githubusercontent.com/TheAlgorithms/Python/master/searches/  
    ↪double_linear_search_recursion.py",  
    "https://raw.githubusercontent.com/KosingZhu/tensorflow/master/tensorflow/  
    ↪python/tools/module_util.py",  
    "https://raw.githubusercontent.com/EricRemmerswaal/tensorflow/master/  
    ↪tensorflow/python/distribute/distribute_coordinator_context.py",  
    "https://raw.githubusercontent.com/computationalartist/tensorflow/master/  
    ↪tensorflow/python/ops/numpy_ops/integration_test/benchmarks/numpy_mlp.py",  
    "https://raw.githubusercontent.com/Van-an/tensorflow/master/tensorflow/  
    ↪python/distribute/coordinator/values.py",  
    "https://raw.githubusercontent.com/nkgwer/tensorflow/master/tensorflow/lite/  
    ↪tools/visualize.py",  
]
```

```

    "https://raw.githubusercontent.com/gitblazer/youtube-dl/master/youtube_dl/
↪version.py",
    "https://raw.githubusercontent.com/Joshua-Barawa/My-Photos/master/venv/lib/
↪python3.8/site-packages/django/contrib/messages/__init__.py",
    "https://raw.githubusercontent.com/PaliC/pytorch/master/test/fx/
↪test_subgraph_rewriter.py"
]

```

Retrieve the python scripts:

```

[ ]: for url in urls:
    print(f"Working on url: {url}")
    response = requests.get(url)
    file_name = os.path.basename(url)
    file_path = os.path.join(code_dir, file_name)

    with open(file_path, "wb") as file:
        file.write(response.content)

```

```

[ ]: files = os.listdir(code_dir)
    for file in files:
        print(file)

```

Concatenate scripts into a list:

```

[ ]: code_dataset = []
    for file in os.listdir(code_dir):
        code_dataset.append(
            {'text': open(os.path.join(code_dir, file), 'r').read()}
        )

```

Convert list to Hugging Face Dataset object:

```

[ ]: code_dataset = datasets.Dataset.from_list(code_dataset)
    print(code_dataset)

```

Combine the python code dataset with the pretraining dataset you downloaded above:

```

[ ]: dataset = datasets.concatenate_datasets(
    [pretraining_dataset, code_dataset]
)
    print(dataset)

```

1.2 2. Data cleaning

In the cells below, you'll carry out the following cleaning steps: 1. Filter out samples that are too short 2. Remove repetitions within a single text example 3. Remove duplicated documents 4. Quality filter to remove non-English texts


```
[ ]: dataset.num_rows
```

1.2.1 Remove examples that are too short

```
[ ]: import heapq

def paragraph_length_filter(x):
    """Returns False iff a page has too few lines or lines are too short."""
    lines = x['text'].split('\n')
    if (
        len(lines) < 3
        or min(heapq.nlargest(3, [len(line) for line in lines])) < 3
    ):
        return False
    return True
```

```
[ ]: dataset = dataset.filter(
    paragraph_length_filter,
    load_from_cache_file=False
)
```

```
[ ]: dataset.num_rows
```

1.2.2 Remove repeated text within training examples

Here you'll remove text repetitions within each example.

```
[ ]: def find_duplicates(paragraphs):
    """
    Use this function to find the number of repetitions
    in the paragraphs.
    """
    unique_x = set()
    duplicate_chars = 0
    duplicate_elements = 0
    for element in paragraphs:
        if element in unique_x:
            duplicate_chars += len(element)
            duplicate_elements += 1
        else:
            unique_x.add(element)
    return duplicate_elements, duplicate_chars
```

```
[ ]: import re

def paragraph_repetition_filter(x):
    """
    Returns False iff a page has too many repetitions.
```

```

"""
text = x['text']
paragraphs = re.compile(r"\n{2,}").split(text.strip()) #
↳ Split by paragraphs (2 or more newlines)
paragraphs_duplicates, char_duplicates = find_duplicates(paragraphs) #
↳ Find number of duplicates in paragraphs
if paragraphs_duplicates / len(paragraphs) > 0.3:
    return False
if char_duplicates / len(text) > 0.2:
    return False
return True

```

```
[ ]: dataset = dataset.filter(
    paragraph_repetition_filter,
    load_from_cache_file=False
)
```

```
[ ]: dataset.num_rows
```

1.2.3 Deduplication

In this section, you'll remove duplicate examples from the entire dataset (in contrast to the previous step where you were just looking for repeated text in each example.)

```
[ ]: def deduplication(ds):
    def dedup_func(x):
        """Use this function to remove duplicate entries"""
        if x['text'] in unique_text:
            return False
        else:
            unique_text.add(x['text'])
            return True

    unique_text = set()

    ds = ds.filter(dedup_func, load_from_cache_file=False, num_proc=1)
    return ds

dataset = deduplication(dataset)
```

```
[ ]: dataset.num_rows
```

1.2.4 Quality filter - Language

Here you'll remove any text examples that are in a language other than English. The code here uses a language detection model called fastText. You can read about fastText [here](#).

```
[ ]: # !pip install fasttext
```

```
[ ]: import urllib
from fasttext.FastText import _FastText

def english_language_filter(ds):
    # load language detection model
    model = _FastText('./models/upstage/L2_language_model.bin')

    def is_english(x):
        # Predict language of the text and probability
        language, score = model.predict(x['text'].replace("\n", ""))

        language = language[0].split("__")[2]
        return score > 0.4 and language == "en" # change code here if building
        ↪ a model in another language

    ds = ds.filter(is_english, load_from_cache_file=False, num_proc=1)
    return ds

dataset = english_language_filter(dataset)

[ ]: dataset.num_rows
```

1.3 3. Save the dataset to disk

Read more about the parquet data format [here](#).

```
[ ]: file_path = "./data/preprocessed_dataset.parquet"
dataset.to_parquet(file_path)
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

Lesson 3: Data Packaging

1. Tokenizing and creating input_ids

Start by loading the dataset from the previous lesson:

```
In [ ]: import datasets

dataset = datasets.load_dataset(
    "parquet",
    data_files="./data/preprocessed_dataset.parquet",
    split="train"
)
print(dataset)
```

Use the `shard` method of the Hugging Face `Dataset` object to split the dataset into 10 smaller pieces, or *shards* (think shards of broken glass). You can read more about sharding at [this link](#).

```
In [ ]: dataset = dataset.shard(num_shards=10, index=0)
print(dataset)
```

Load the tokenizer and try it out:

```
In [ ]: from transformers import AutoTokenizer
model_path_or_name = "./models/SOLAR-10.7B-v1.0"
tokenizer = AutoTokenizer.from_pretrained(
    model_path_or_name,
    use_fast=False
)
```

```
In [ ]: tokenizer.tokenize("I'm a short sentence")
```

Create a helper function:

```
In [ ]: def tokenization(example):
    # Tokenize
    tokens = tokenizer.tokenize(example["text"])

    # Convert tokens to ids
    token_ids = tokenizer.convert_tokens_to_ids(tokens)

    # Add <bos>, <eos> tokens to the front and back of tokens_ids
    # bos: begin of sequence, eos: end of sequence
    token_ids = [
        tokenizer.bos_token_id] \
        + token_ids \
        + [tokenizer.eos_token_id]
    example["input_ids"] = token_ids

    # We will be using this column to count the total number of tokens
    # in the final dataset
    example["num_tokens"] = len(token_ids)
    return example
```

Tokenize all the examples in the pretraining dataset:

```
In [ ]: dataset = dataset.map(tokenization, load_from_cache_file=False)
print(dataset)
```

```
In [ ]: sample = dataset[3]

print("text", sample["text"][:30]) #
print("\ninput_ids", sample["input_ids"][:30])
print("\nnum_tokens", sample["num_tokens"])
```

Check the total number of tokens in the dataset:

```
In [ ]: import numpy as np
np.sum(dataset["num_tokens"])
```

2. Packing the data



Packing data for training

Concatenate input_ids for all examples into a single list:

```
In [ ]: input_ids = np.concatenate(dataset["input_ids"])
print(len(input_ids))
```

```
In [ ]: max_seq_length = 32
```

```
In [ ]: total_length = len(input_ids) - len(input_ids) % max_seq_length
print(total_length)
```

Discard extra tokens from end of the list so number of tokens is exactly divisible by `max_seq_length` :

```
In [ ]: input_ids = input_ids[:total_length]
print(input_ids.shape)
```

```
In [ ]: input_ids_resaped = input_ids.reshape(-1, max_seq_length).astype(np.int32)
input_ids_resaped.shape
```

```
In [ ]: type(input_ids_resaped)
```

Convert to Hugging Face dataset:

```
In [ ]: input_ids_list = input_ids_resaped.tolist()
packaged_pretrain_dataset = datasets.Dataset.from_dict(
    {"input_ids": input_ids_list}
)
print(packaged_pretrain_dataset)
```

3. Save the packed dataset to disk

```
In [ ]: packaged_pretrain_dataset.to_parquet("./data/packaged_pretrain_dataset.parquet")
```


Lesson_4

March 5, 2025

1 Lesson 4: Preparing your model for training

```
[ ]: # Ignore insignificant warnings (ex: deprecation warnings)
import warnings
warnings.filterwarnings('ignore')

# Set a seed value for reproducibility
import torch

def fix_torch_seed(seed=42):
    torch.manual_seed(seed)
    torch.cuda.manual_seed(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

fix_torch_seed()
```

1.1 1. Model configuration

You'll configure models based on Meta's Llama family of models. The transformers library has several tools for working with these models, which you can read about [here](#).

Start by creating a LlamaConfig object to configure the architecture of the model:

```
[ ]: from transformers import LlamaConfig
config = LlamaConfig()
print(config)
```

Next, update parameters to change the model architecture:

```
[ ]: config.num_hidden_layers = 12          # reduced from 32 to 12
config.hidden_size = 1024                 # reduced 1/4 from 4096 to 1024
config.intermediate_size = 4096          # reduced 1/3 from 11008 to 4096 (dimension
    ↳ of MLP representations)
config.num_key_value_heads = 8            # reduced 1/4 from 32 to 8 (defaults to
    ↳ num_attention_heads=32)
config.torch_dtype = "bfloat16"          # for half-precision training
config.use_cache = False                 # `True` is incompatible w/ gradient
    ↳ checkpointing
```

```
print(config)
```

1.2 2. Weight initialization

In the next sections, you'll explore four different ways to initialize the weights of a model for training: 1. Random weight initialization 2. Using an existing model for continued pre-training 3. Downscaling an existing model 4. Upscaling an existing model

1.2.1 Random weight initialization

Randomly initializing model weights sets all weights to values from a truncated normal distribution with mean 0 and standard deviation of 0.02. Values beyond 2-sigma from the mean are set to 0.

```
[ ]: from transformers import LlamaForCausalLM
model = LlamaForCausalLM(config)
print(model)

[ ]: def print_nparams(model):
    """Calculate the total number of model parameters"""
    nparams = sum(p.numel() for p in model.parameters())
    print(f"The total number of parameters is: {nparams}")

print_nparams(model)  # 248013824 => 248M
```

Take a look at a sample of the weights in a single layer:

```
[ ]: layer_name = "model.layers.0.self_attn.q_proj.weight"

for name, param in model.named_parameters():
    if name == layer_name:
        print(f"First 30 weights of layer '{layer_name}':")
        print(param.data.view(-1)[:30])
        break
```

Try using the model for inference:

```
[ ]: # Load a tokenizer from Upstage Solar,
# which is compatible with the Llama-2 tokenizer
from transformers import LlamaTokenizer
model_dir = "./models/SOLAR-10.7B-v1.0"
tokenizer = LlamaTokenizer.from_pretrained(model_dir)

# Run simple inference with prompt
from transformers import TextStreamer

prompt = "I am an engineer. I love"

inputs = tokenizer(prompt, return_tensors="pt").to(model.device)
```



```

streamer = TextStreamer(
    tokenizer,
    skip_prompt=True,
    skip_special_tokens=True
)

outputs = model.generate(
    **inputs,
    streamer=streamer,
    use_cache=True,
    max_new_tokens=128,
    do_sample=False
)

```

Remove the model from memory to avoid crashing the kernel:

```

[ ]: # NOTE: We're running large models in a limited environment. Run me if you
    ↪ encounter any memory issues.
import gc
del model
del streamer
del outputs
gc.collect()

```

1.2.2 Reuse general pretrained model weights

If you load an existing model, you can use it as is to continue pretraining on new data.

```

[ ]: from transformers import AutoModelForCausalLM

model_name_or_path = "./models/TinySolar-248m-4k"
model = AutoModelForCausalLM.from_pretrained(
    model_name_or_path,
    device_map="cpu",
    torch_dtype=torch.bfloat16,
)

```

Remove the model from memory to avoid crashing the kernel:

```

[ ]: # NOTE: We're running large models in a limited environment. Run me if you
    ↪ encounter any memory issues.
del model
gc.collect()

```

1.2.3 Downscaling from a general pretrained model

Here you'll downscale the tinySolar-248m-4k model from a 12 layer model to a 10 layer model.

```
[ ]: from transformers import AutoTokenizer, AutoConfig

model_name_or_path = "./models/TinySolar-248m-4k"
model = AutoModelForCausalLM.from_pretrained(
    model_name_or_path,
    device_map="cpu",
    torch_dtype=torch.bfloat16,
)
tokenizer = AutoTokenizer.from_pretrained(model_name_or_path)
```

```
[ ]: print(model)
```

```
[ ]: print_nparams(model) # 248013824 => 248M
```

Remove the middle two layers (layers 5 and 6) and update the configuration:

```
[ ]: layers = model.model.layers
model.model.layers = layers[:5] + layers[-5:]

config = AutoConfig.from_pretrained(
    model_name_or_path,
    num_hidden_layers=len(model.model.layers),
)
model.config = config

print_nparams(model) # 217601024 => 217M
```

Clear the memory to avoid crashing the kernel:

```
[ ]: # NOTE: We're running large models in a limited environment. Run me if you
    ↪ encounter any memory issues.
import gc
del model
gc.collect()
```

1.2.4 Depth Upscaling from a general pretrained model

Here you are going to upscale the tinySolar-248m-4k model from 12 layers to 16 layers. Here are the steps you'll take: 1. Configure a 16 layer model and initialize it with random weights 2. Load the 12 layer tinySolar-248m-4k model into memory 3. Copy the bottom 8 and top 8 layers from the 12 layer model and use them to overwrite the random weights of the 16 layer model 4. Copy over the embedding and classifying layers to replace the randomly initialized counterparts in the 16 layer model

```
[ ]: config = LlamaConfig(
    num_hidden_layers=16, # We want our model to have 16 final layers
    hidden_size=1024,
    intermediate_size=4096,
    num_attention_heads=32,
```

```

    num_key_value_heads=8,
    torch_dtype="bfloat16",
    use_cache=False
)
print(config)

```

```

[ ]: model = LlamaForCausalLM(config)
model = model.to(dtype=torch.bfloat16) # convert to bfloat16
print_nparams(model) # 308839424 => 308M

```

```

[ ]: model_name_or_path = "upstage/TinySolar-248m-4k"
pretrained_model = AutoModelForCausalLM.from_pretrained(
    model_name_or_path,
    device_map="cpu",
    torch_dtype=torch.bfloat16,
)
tokenizer = AutoTokenizer.from_pretrained(model_name_or_path)

print_nparams(pretrained_model) # 248013824 => 248M

```

```

[ ]: from copy import deepcopy

model.model.layers = deepcopy(pretrained_model.model.layers[:-4]) \
    + deepcopy(pretrained_model.model.layers[4:])

model.model.embed_tokens = deepcopy(pretrained_model.model.embed_tokens)

model.lm_head = deepcopy(pretrained_model.lm_head)

print(model.config)

```

Check the number of parameters is still 308 million:

```

[ ]: print_nparams(model) # 308839424 => 308M

```

Try using the model for inference:

```

[ ]: # Run simple inference to show no trained model
prompt = "I am an engineer. I love"

inputs = tokenizer(prompt, return_tensors="pt").to(model.device)

streamer = TextStreamer(
    tokenizer,
    skip_prompt=True,
    skip_special_tokens=True
)

```

```
outputs = model.generate(  
    **inputs,  
    streamer=streamer,  
    use_cache=True,  
    max_new_tokens=128,  
    do_sample=False  
)
```

1.2.5 Save the model to disk

Note the new model name here which reflects the 308 million parameters of the new, upscaled model.

```
[ ]: model.save_pretrained('./data/TinySolar-308m-4k-init')
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

Lesson_5

March 5, 2025

1 Lesson 5. Model training

Pretraining is very expensive! Please check costs carefully before starting a pretraining project.

You can get a rough estimate your training job cost using [this calculator](#) from Hugging Face. For training on other infrastructure, e.g. AWS or Google Cloud, please consult those providers for up to date cost estimates.

```
[ ]: import warnings
warnings.filterwarnings('ignore')
```

1.1 1. Load the model to be trained

Load the upscaled model from the previous lesson:

```
[ ]: import torch
from transformers import AutoModelForCausalLM

pretrained_model = AutoModelForCausalLM.from_pretrained(
    "./models/TinySolar-308m-4k-init",
    device_map="cpu",
    torch_dtype=torch.bfloat16,
    use_cache=False,
)
```

```
[ ]: pretrained_model
```

1.2 2. Load dataset

Here you'll update two methods on the `Dataset` object to allow it to interface with the trainer. These will be applied when you specify the dataset you created in Lesson 3 as the training data in the next section.

Note that the code has additional comment strings that don't appear in the video. These are to help you understand what each part of the code is doing.

```
[ ]: import datasets
from torch.utils.data import Dataset

class CustomDataset(Dataset):
```

```

def __init__(self, args, split="train"):
    """Initializes the custom dataset object."""
    self.args = args
    self.dataset = datasets.load_dataset(
        "parquet",
        data_files=args.dataset_name,
        split=split
    )

def __len__(self):
    """Returns the number of samples in the dataset."""
    return len(self.dataset)

def __getitem__(self, idx):
    """
    Retrieves a single data sample from the dataset
    at the specified index
    """
    # Convert the lists to a LongTensor for PyTorch
    input_ids = torch.LongTensor(self.dataset[idx]["input_ids"])
    labels = torch.LongTensor(self.dataset[idx]["input_ids"])

    # Return the sample as a dictionary
    return {"input_ids": input_ids, "labels": labels}

```

1.3 3. Configure Training Arguments

Here you set up the training run. The training dataset you created in Lesson 3 is specified in the Dataset configuration section.

Note: there are comment strings in the cell below that don't appear in the video. These have been included to help you understand what each parameter does.

```

[ ]: from dataclasses import dataclass, field
import transformers

@dataclass
class CustomArguments(transformers.TrainingArguments):
    dataset_name: str = field(
        default="./parquet/packaged_pretrain_dataset.parquet")
        # Dataset configuration
    num_proc: int = field(default=1)
        # Number of
    ↪subprocesses for data preprocessing
    max_seq_length: int = field(default=32)
        # Maximum sequence
    ↪length

    # Core training configurations
    seed: int = field(default=0)
        # Random seed for
    ↪initialization, ensuring reproducibility

```

```

    optim: str = field(default="adamw_torch")           # Optimizer, here it's
    ↪AdamW implemented in PyTorch
    max_steps: int = field(default=30)                 # Number of maximum
    ↪training steps
    per_device_train_batch_size: int = field(default=2) # Batch size per
    ↪device during training

    # Other training configurations
    learning_rate: float = field(default=5e-5)         # Initial learning
    ↪rate for the optimizer
    weight_decay: float = field(default=0)             # Weight decay
    warmup_steps: int = field(default=10)              # Number of steps for
    ↪the learning rate warmup phase
    lr_scheduler_type: str = field(default="linear")    # Type of learning
    ↪rate scheduler
    gradient_checkpointing: bool = field(default=True)  # Enable gradient
    ↪checkpointing to save memory
    dataloader_num_workers: int = field(default=2)     # Number of
    ↪subprocesses for data loading
    bf16: bool = field(default=True)                  # Use bfloat16
    ↪precision for training on supported hardware
    gradient_accumulation_steps: int = field(default=1) # Number of steps to
    ↪accumulate gradients before updating model weights

    # Logging configuration
    logging_steps: int = field(default=3)              # Frequency of logging
    ↪training information
    report_to: str = field(default="none")             # Destination for
    ↪logging (e.g., WandB, TensorBoard)

    # Saving configuration
    # save_strategy: str = field(default="steps")       # Can be replaced
    ↪with "epoch"
    # save_steps: int = field(default=3)               # Frequency of
    ↪saving training checkpoint
    # save_total_limit: int = field(default=2)          # The total number
    ↪of checkpoints to be saved

```

Parse the custom arguments and set the output directory where the model will be saved:

```

[ ]: parser = transformers.HfArgumentParser(CustomArguments)
args, = parser.parse_args_into_dataclasses(
    args=["--output_dir", "output"]
)

```

Setup the training dataset:

```
[ ]: train_dataset = CustomDataset(args=args)
```

Check the shape of the dataset:

```
[ ]: print("Input shape: ", train_dataset[0]['input_ids'].shape)
```

1.4 4. Run the trainer and monitor the loss

First, set up a callback to log the loss values during training (note this cell is not shown in the video):

```
[ ]: from transformers import Trainer, TrainingArguments, TrainerCallback

# Define a custom callback to log the loss values
class LossLoggingCallback(TrainerCallback):
    def on_log(self, args, state, control, logs=None, **kwargs):
        if logs is not None:
            self.logs.append(logs)

    def __init__(self):
        self.logs = []

# Initialize the callback
loss_logging_callback = LossLoggingCallback()
```

Then, create an instance of the Hugging Face Trainer object from the `transformers` library. Call the `train()` method of the trainer to initialize the training run:

```
[ ]: from transformers import Trainer

trainer = Trainer(
    model=pretrained_model,
    args=args,
    train_dataset=train_dataset,
    eval_dataset=None,
    callbacks=[loss_logging_callback]
)

trainer.train()
```

You can use the code below to save intermediate model checkpoints in your own training run:

```
[ ]: # Saving configuration
      # save_strategy: str = field(default="steps")           # Can be replaced
      ↪with "epoch"                                           # Frequency of
      # save_steps: int = field(default=3)                   ↪saving training checkpoint
      # save_total_limit: int = field(default=2)             # The total number
      ↪of checkpoints to be saved
```


1.4.1 Checking the performance of an intermediate checkpoint

Below, you can try generating text using an intermediate checkpoint of the model. This checkpoint was saved after 10,000 training steps. As you did in previous lessons, you'll use the Solar tokenizer and then set up a `TextStreamer` object to display the text as it is generated:

```
[ ]: from transformers import AutoTokenizer, TextStreamer
model_name_or_path = "./models/TinySolar-248m-4k"
tokenizer = AutoTokenizer.from_pretrained(model_name_or_path)

[ ]: from transformers import AutoTokenizer, TextStreamer, AutoModelForCausalLM
import torch

model_name_or_path = "./models/upstage/output/checkpoint-10000"
model2 = AutoModelForCausalLM.from_pretrained(
    model_name_or_path,
    device_map="auto",
    torch_dtype=torch.bfloat16,
)

[ ]: prompt = "I am an engineer. I love"

inputs = tokenizer(prompt, return_tensors="pt").to(model2.device)

streamer = TextStreamer(
    tokenizer,
    skip_prompt=True,
    skip_special_tokens=True
)

outputs = model2.generate(
    **inputs,
    streamer=streamer,
    use_cache=True,
    max_new_tokens=64,
    do_sample=True,
    temperature=1.0,
)
```

Lesson_6

March 5, 2025

1 Lesson 6. Model evaluation

The model comparison tool that Sung described in the video can be found at this link: <https://console.upstage.ai/> (note that you need to create a free account to try it out.)

A useful tool for evaluating LLMs is the **LM Evaluation Harness** built by EleutherAI. Information about the harness can be found at this [github repo](#):

You can run the commented code below to install the evaluation harness in your own environment:

```
[ ]: #!pip install -U git+https://github.com/EleutherAI/lm-evaluation-harness
```

You will evaluate TinySolar-248m-4k on 5 questions from the **TruthfulQA MC2 task**. This is a multiple-choice question answering task that tests the model's ability to identify true statements. You can read more about the TruthfulQA benchmark in [this paper](#), and you can checkout the code for implementing the tasks at this [github repo](#).

The code below runs only the TruthfulQA MC2 task using the LM Evaluation Harness:

```
[ ]: !lm_eval --model hf \
    --model_args pretrained=./models/TinySolar-248m-4k \
    --tasks truthfulqa_mc2 \
    --device cpu \
    --limit 5
```

1.0.1 Evaluation for the Hugging Face Leaderboard

You can use the code below to test your own model against the evaluations required for the [Hugging Face leaderboard](#).

If you decide to run this evaluation on your own model, don't change the few-shot numbers below - they are set by the rules of the leaderboard.

```
[ ]: import os

def h6_open_llm_leaderboard(model_name):
    task_and_shot = [
        ('arc_challenge', 25),
        ('hellaswag', 10),
        ('mmlu', 5),
        ('truthfulqa_mc2', 0),
```

```

        ('winogrande', 5),
        ('gsm8k', 5)
    ]

    for task, fewshot in task_and_shot:
        eval_cmd = f"""
        lm_eval --model hf \
            --model_args pretrained={model_name} \
            --tasks {task} \
            --device cpu \
            --num_fewshot {fewshot}
        """
        os.system(eval_cmd)

h6_open_llm_leaderboard(model_name="YOUR_MODEL")

```